

# CS5785\_HW1\_Writeup

September 13, 2017

## 1 Digit Recognizer

Group Members: Huajun Bai(hb364), Hao Zheng(hz466)

### 1.1 Q1a Loading Data

```
In [2]: # Importing library
import collections
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn.metrics.pairwise as skl
from scipy import stats
from scipy.optimize import brentq
from scipy.interpolate import interp1d
from sklearn import linear_model, cross_validation, preprocessing, metrics



In [3]: # load data
data = np.loadtxt(fname = 'train.csv', delimiter = ',', skiprows=1)
n,p = data.shape
```



### 1.2 Q2b



Below is the function to display any digit d. And a display of digits 0-9 in the data.



```
In [4]: # A function to display a digit d
def display(data, d):
    for i in range(n):
        if data[i, 0] == d:
            plt.matshow(data[i,1:785].reshape(28,28), cmap='gray')
            plt.show()
            break;



In [6]: for j in range(10):
        display(data, j)
```












### 1.3 Q1c

Below is data for number of each digit in the data set. The distribution is not uniform, so the graph is not even.

```
In [165]: collections.Counter(data[:,0])
```

```
Out[165]: Counter({0.0: 4132,  
                  1.0: 4684,  
                  2.0: 4177,  
                  3.0: 4351,  
                  4.0: 4072,  
                  5.0: 3795,  
                  6.0: 4137,  
                  7.0: 4401,  
                  8.0: 4063,  
                  9.0: 4188})
```

```
In [90]: plt.hist(data[:,0], normed = True, histtype = 'bar')  
plt.xticks(np.arange(10), ('0','1','2','3','4','5','6','7','8','9'))  
plt.show()
```



### 1.4 Q1d

Below you can find code and graphs for each digit and their nearest neighbor. No.3 is a outlier, where nearest neighbor outputs a '5' as its nearest neighbor.


```

In [11]: def findNearest(index):
          diff = np.subtract(np.array(data[:, 1:785]), np.array(data[index, 1:785]));
          distM = np.square(diff) # delete index itself
          # Assume no same data set exist, i.e. no 0 dist except itself
          return np.argsort(np.sum(distM, axis=1)).item(1)

In [12]: # Part d
          # No.3 outlier
          for i in range(10):
              for j in range(n):
                  if(data[j,0] == i):
                      plt.matshow(data[j,1:785].reshape(28,28), cmap='gray')
                      print ("this digit is " + repr(i)+ " from No." + repr(j));
                      plt.show()
                      nearIndex = findNearest(j)
                      print ("Nearest digit is "+ repr(data[nearIndex, 0])+ " from No." + repr(ne
                      plt.matshow(data[nearIndex,1:785].reshape(28,28), cmap='gray')
                      plt.show()
                      break


```

this digit is 0 from No.1




Nearest digit is 0.0 from No.12950







this digit is 1 from No.0




Nearest digit is 1.0 from No.29704




this digit is 2 from No.16




Nearest digit is 2.0 from No.9536




this digit is 3 from No.7




Nearest digit is 5.0 from No.8981




this digit is 4 from No.3




Nearest digit is 4.0 from No.14787




this digit is 5 from No.8




Nearest digit is 5.0 from No.30073




this digit is 6 from No.21



Nearest digit is 6.0 from No.16240




this digit is 7 from No.6







Nearest digit is 7.0 from No.15275




this digit is 8 from No.10




Nearest digit is 8.0 from No.32586



this digit is 9 from No.11



Nearest digit is 9.0 from No.35742




## 1.5 Q1e

Computing pairwise distances of all 0's and 1's. Plotting histogram of genuine (00 and 11 distances combined) and imposter (01 distances) distances. Genuine distances are marked green and imposter distances are marked red below.

```
In [18]: zeros = data[data[:,0]==0]
ones = data[data[:,0]==1]
pairDist00 = skl.pairwise_distances(zeros)
pairDist11 = skl.pairwise_distances(ones)
pairDist01 = skl.pairwise_distances(zeros, ones)
pairGE = np.append(pairDist00, pairDist11)

In [20]: plt.hist(pairGE.flatten(), bins = 'auto', normed = True, rwidth = 0.5, color = 'g')
#plt.hist(pairDist11.flatten(), bins = 'auto', normed = True, rwidth = 0.3, color = 'g')
plt.hist(pairDist01.flatten(), bins = 'auto', normed = True, rwidth = 0.5, color = 'r')
plt.show()
```



## 1.6 Q1f


Generating ROC curve below with true positive rate (tpr) as x-axis and false positive rate (fpr) as y-axis. We find equal error rate is 0.18554865605818877 for this training data. If we just guess randomly, i.e. whenever there is a test input, we have 50% chance to guess it as a "0" and 50%

chance to guess it as a "1". Then the expectation of error rate will be ratio of 0's vs 1's in the test set.

```
In [23]: fpr = [] # false_positive_rate
         tpr = [] # true_positive_rate
         lenGE = float(len(pairGE))
         len01 = float(len(pairDist01.flatten()))
         for thr in range(0, 4500, 5):
             tmp1 = float(np.count_nonzero(pairGE < thr))/lenGE
             tmp2 = float(float(np.count_nonzero(pairDist01<thr))/len01)
             tpr.append(tmp1)
             fpr.append(tmp2)

In [24]: # This is the ROC curve
         plt.plot(fpr, tpr)
         plt.xlabel("tpr")
         plt.ylabel("fpr")
         # This is the AUC
         auc = np.trapz(tpr,fpr)

         x = np.linspace(0.0, 1.0, num=100)
         y = 1-x
         plt.plot(y,x)
         plt.show()
```



```
In [25]: eer = brentq(lambda x : 1. - x - interp1d(fpr, tpr)(x), 0., 1.)
        print ("EER is " + repr(eer))
```

EER is 0.18554865605818877

## 1.7 Q1g

Below is our implementation of kNN, which is able to take in a test set of dimension  $m \times p$ , i.e.  $m$  data points each with  $p$  features and output labels for each of  $m$  points.

```
In [27]: #implement kNN
        def kNN(data, label, test, k):
            # ndarray data has dimension n*p
            # label has dimension n
            # ndarray test has dimension m*p
            # y has dimension m
            y = []
            Dist = skl.pairwise_distances(data, test) # n*m
            #idx = np.argpartition(Dist, k, axis = 0)
            #ypredict = label[idx, np.arange(Dist.shape[1])[None,:]] #
            #return stats.mode(ypredict[:k], axis = )
            for i in range(test.shape[0]):
                idx = np.argpartition(Dist[:,i], k)
                y.append(stats.mode(label[idx[:k]]).mode[0])
            return y
```

## 1.8 Q1h

Here is 3-fold analysis on training data with  $k=1$ . The average accuracy is 0.9643809523809667.

```
In [29]: k = 1
        matrixList = []
        error = []
        for trainKF, testKF in (cross_validation.KFold(len(data), n_folds=3)):
            predicted = kNN(data[trainKF,1:785], data[trainKF, 0], data[testKF, 1:785], k)
            error.append(np.mean(predicted != data[testKF, 0]))
            matrixList.append(metrics.confusion_matrix(data[testKF, 0], predicted))
```

```
In [30]: print ("Average accuracy on k=1, 3-fold is: " + repr(1-np.mean(error)))
```

Average accuracy on k=1, 3-fold is: 0.96438095238095234

## 1.9 Q1i

Display the confusion matrix of 3-fold computation below. Number 8 is the most tricky to classify.

```
In [46]: conf = np.zeros((10,10), dtype=np.int)
        for m in matrixList:
            conf = np.add(conf, m)
        print("Here is the confusion matrix:")
        print(conf)
        np.fill_diagonal(conf, 0)
        missC = np.sum(conf, axis = 1)
        print ("\n Number of misclassification for each number below:")
        print(missC)
        print("\n The most tricky number to predict is " + repr(np.argmax(missC)))
```

Here is the confusion matrix:

```
[[4099   1   5   1   0   6  16   0   2   2]
 [  0 4646   9   3   5   2   4   9   4   2]
 [ 29  24 4000  21   5   5   5  70  13   5]
 [  2   8  31 4148   0  75   1  26  37  23]
 [  2  39   0   0 3892   1  12  14   1 111]
 [  9   5   1  61   6 3622  48   4  14  25]
 [ 30   7   0   1   6  19 4073   0   1   0]
 [  1  43  15   2  12   0   0 4264   0  64]
 [ 12  33  18  74  13  59  16  12 3782  44]
 [ 12   9   3  25  60  13   3  73  12 3978]]
```

Number of misclassification for each number below:  
[ 33 38 177 203 180 173 64 137 281 210]

The most tricky number to predict is 8

## 1.10 Q1j

Import test data and generate predictions for test data with my kNN. Submission screenshot included.

```
In [47]: test = np.loadtxt(fname = 'test.csv', delimiter = ',', skiprows = 1)
```

```
In [52]: results = kNN(data[:,1:785], data[:, 0], test, 1)
```

```
In [53]: df = pd.DataFrame(results)
        df.index.name='ImageId'
        df.index+=1
        df.columns=['Label']
        df.to_csv('results.csv', header=True)
```

```
In [54]: from IPython.display import Image
        Image("digits.png")
```

Out[54]:

All   Successful   Selected		
Submission and Description	Public Score	Use for Final Score
<a href="#">results.csv</a> 4 days ago by <a href="#">HuajunBai</a> CS5785 HW1 submitting	0.97114	<input type="checkbox"/>



# myTitanic

September 13, 2017

## 1 The Titanic Disaster

### 1.1 Q2a

Importing libraries and data below

```
In [1]: # import libraries
import pandas as pd
import numpy as np
from sklearn import preprocessing, cross_validation, datasets, neighbors, linear_model

C:\Users\Huajun\Anaconda3\lib\site-packages\sklearn\cross_validation.py:44: DeprecationWarning:
  "This module will be removed in 0.20.", DeprecationWarning)

In [2]: # import data
train = pd.read_csv("train.csv", dtype={"Age": np.float64},)
test = pd.read_csv("test.csv", dtype={"Age": np.float64},)

In [3]: #Print to standard output, and see the results in the "log" section below after running
print("\n\nTop of the training data:")
print(train.head())
```

Top of the training data:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	

4			Allen, Mr. William Henry	male	35.0	0
---	--	--	--------------------------	------	------	---

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

## 1.2 Q2b

Choose 3 features: "Sex", "Age" and "PClass" according to [this article](#) Then perform 5-fold cross-validation using logistic regression. Expected score (75%, 80%).

```
In [21]: # Fill missing age values for the train and test data with corresponding mean value,
# and convert values from float to integer.
train.loc[train["Sex"] == "male", "Sex"] = 0
train.loc[train["Sex"] == "female", "Sex"] = 1

test.loc[test["Sex"] == "male", "Sex"] = 0
test.loc[test["Sex"] == "female", "Sex"] = 1

train["Age"] = train["Age"].fillna(train["Age"].mean())
train['Age'] = train['Age'].astype(int)

test["Age"] = test["Age"].fillna(test["Age"].mean())
test['Age'] = test['Age'].astype(int)

In [39]: # choose features
# pclass--2, sex--4, age--5
data = pd.DataFrame.as_matrix(train)
myData = data[:, [0,1,2,4,5]]

In [31]: logistic = linear_model.LogisticRegression()
for trainKF, testKF in (cross_validation.KFold(len(myData), n_folds=5)):
    print('LogisticRegression score: %f'
          % logistic.fit(myData[trainKF, 2:5], list(myData[trainKF, 1])).score(myData[t
```

LogisticRegression score: 0.798883  
LogisticRegression score: 0.814607  
LogisticRegression score: 0.775281  
LogisticRegression score: 0.752809  
LogisticRegression score: 0.808989

## 1.3 Q2c

Computing test prediction and submitting below. Accuracy achieved at 75.6%

```
In [10]: test.head()
```

```
Out[10]:
```

	PassengerId	Pclass	Name	Sex
0	892	3	Kelly, Mr. James	male
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female
2	894	2	Myles, Mr. Thomas Francis	male
3	895	3	Wirz, Mr. Albert	male
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	34	0	0	330911	7.8292	NaN	Q
1	47	1	0	363272	7.0000	NaN	S
2	62	0	0	240276	9.6875	NaN	Q
3	27	0	0	315154	8.6625	NaN	S
4	22	1	1	3101298	12.2875	NaN	S

```
In [35]: results = logistic.predict(pd.DataFrame.as_matrix(test)[:,[1,3,4]])
```


```
In [37]: #-----RESULT SUBMISSION-----#
submission = pd.DataFrame({
    "PassengerId": test["PassengerId"],
    "Survived": results
})

submission.to_csv('pred.csv', index=False)
```

```
In [38]: from IPython.display import Image
Image("titanic.png")
```

Out[38]:

6833newHuajunBai

0.7559811m

Your Best Entry ↑

Your submission scored 0.75598, which is not an improvement of your best score. Keep trying!

# **CS 5785 Applied Machine Learning**

## **Homework 1**

**Huajun Bai / hb364**

**Hao Zheng / hz466**

## Written Exercise

Q1

Rule 1:  $Var[X] = E[X^2] - E[X]^2$

Rule 2:  $Cov[X, Y] = E[XY] - E[X]E[Y]$

$$\begin{aligned} Var(X - Y) &= E[(X - Y)^2] - E[X - Y]^2 = E[X^2 - 2XY + Y^2] - (E[X] - E[Y])^2 \\ &= E[X^2] - 2E[XY] + E[Y^2] - E[X]^2 + 2E[X]E[Y] - E[Y]^2 \\ &= (E[X^2] - E[X]^2) - 2(E[XY] - E[X]E[Y]) + (E[Y^2] - E[Y]^2) \\ &= Var[X] - 2Cov[X, Y] + Var[Y] \end{aligned}$$

Q2

Let + denotes testing positively and - denotes testing negatively.  $D$  denotes defective and  $N$  denotes not defective. Then

$$P(+|D) = 0.95, P(-|N) = 0.95, P(D) = 1/100000, P(N) = 1 - P(D)$$

a) By Bayes Rule,

$$P(D|+) = \frac{P(+|D)P(D)}{P(+|D)P(D) + P(+|N)P(N)} = 0.019\%$$


b)  $10\text{ m} * P(N)P(+|N) = 99999 * 5 = 499995$

$$10\text{ m} * P(D)P(-|D) = 100 * 0.05 = 5$$

If there are 10 million widgets produced each year, then 499995 good widgets are thrown each year and 5 bad widgets are shipped to customers.

Q3

- a) When  $k=n$ , the classifier just outputs the mode of all data points. Since the data set is chosen as half one class half another class, the prediction error will be 0.5. When  $k=1$ , the classifier will choose the class of the nearest point of the input as an output result. The prediction error will depend on how overlapped the dataset is. In general, when  $k$  decreases from  $n$  to 1, the error should decrease at first, reach its minimum and slight increase at last.
- b) This is two-fold analysis of error rate. When  $k$  decreases from  $n$  to 1, the error should decrease at first, reach its minimum and slight increase at last. Graph looks like this:



- c) How many folds we should use? When number of fold is greater, bias is reduced but variance is increased as each set has smaller number of samples, and computation time is also increased.  
 $K = 10$  usually is not a bad choice. Another way to choose  $K$  is to let  $K$  be a function of sample size, for example set  $K = \sqrt{N}$ .
- d) Since the underlying assumption of kNN is the closer the points are, the more likely they are of the same class. It is intuitive to place higher weights on points that are closer to the test point when  $k$  is large. One way is to set weight for node :  

$$w_v = \frac{1}{\text{dist}(v, x)} / (\sum_{w \in K} \frac{1}{\text{dist}(w, x)})$$
where  $K$  denotes the set of  $k$  nearest nodes.
- e) One reason is that as dimension becomes very high, the distance of points becomes blurred as it is calculated across all features. Distances between points will become insignificant. Another reason is that when dimension is high, the pairwise distance of points will take a lot of computation, thus undesirable.